Logistic Regression

January 8, 2021

[1]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

[2]:
```
ad_data=pd.read_csv('advertising.csv')
```

[3]:
```
ad_data.head()
```

[3]:
```
   Daily Time Spent on Site  Age  Area Income  Daily Internet Usage  \
0                     68.95   35     61833.90                256.09
1                     80.23   31     68441.85                193.77
2                     69.47   26     59785.94                236.50
3                     74.15   29     54806.18                245.89
4                     68.37   35     73889.99                225.58

                          Ad Topic Line            City  Male     Country  \
0     Cloned 5thgeneration orchestration     Wrightburgh     0     Tunisia
1     Monitored national standardization       West Jodi     1       Nauru
2       Organic bottom-line service-desk        Davidton     0  San Marino
3  Triple-buffered reciprocal time-frame  West Terrifurt     1       Italy
4          Robust logistical utilization    South Manuel     0     Iceland

             Timestamp  Clicked on Ad
0  2016-03-27 00:53:11              0
1  2016-04-04 01:39:02              0
2  2016-03-13 20:35:42              0
3  2016-01-10 02:31:19              0
4  2016-06-03 03:36:18              0
```

[4]:
```
ad_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
Daily Time Spent on Site    1000 non-null float64
```

```
Age                       1000 non-null int64
Area Income               1000 non-null float64
Daily Internet Usage      1000 non-null float64
Ad Topic Line             1000 non-null object
City                      1000 non-null object
Male                      1000 non-null int64
Country                   1000 non-null object
Timestamp                 1000 non-null object
Clicked on Ad             1000 non-null int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

[5]:
```
ad_data.describe()
```

[5]:

	Daily Time Spent on Site	Age	Area Income	\
count	1000.000000	1000.000000	1000.000000	
mean	65.000200	36.009000	55000.000080	
std	15.853615	8.785562	13414.634022	
min	32.600000	19.000000	13996.500000	
25%	51.360000	29.000000	47031.802500	
50%	68.215000	35.000000	57012.300000	
75%	78.547500	42.000000	65470.635000	
max	91.430000	61.000000	79484.800000	

	Daily Internet Usage	Male	Clicked on Ad
count	1000.000000	1000.000000	1000.00000
mean	180.000100	0.481000	0.50000
std	43.902339	0.499889	0.50025
min	104.780000	0.000000	0.00000
25%	138.830000	0.000000	0.00000
50%	183.130000	0.000000	0.50000
75%	218.792500	1.000000	1.00000
max	269.960000	1.000000	1.00000

0.0.1 Exploratory Data Analysis

Histogram of the Age

[6]:
```
sns.set_style('whitegrid')
ad_data['Age'].hist(bins=30)
plt.xlabel('Age')
```

[6]:
```
Text(0.5, 0, 'Age')
```

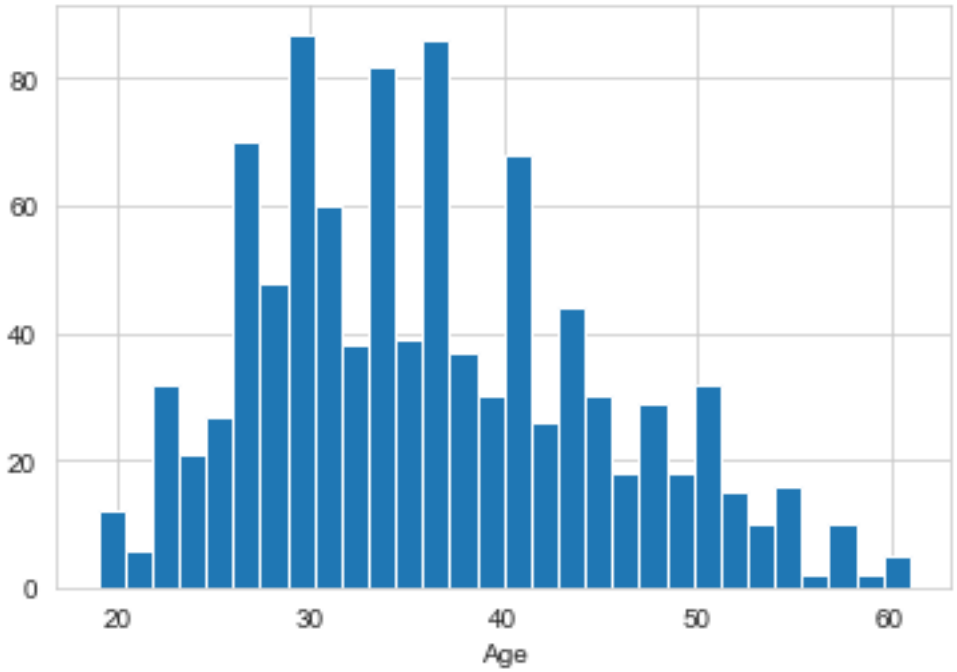

Create a jointplot showing Area Income versus Age.

[7]:
```
sns.jointplot(x='Age',y='Area Income',data=ad_data)
```

[7]:
```
<seaborn.axisgrid.JointGrid at 0x21f58f13808>
```

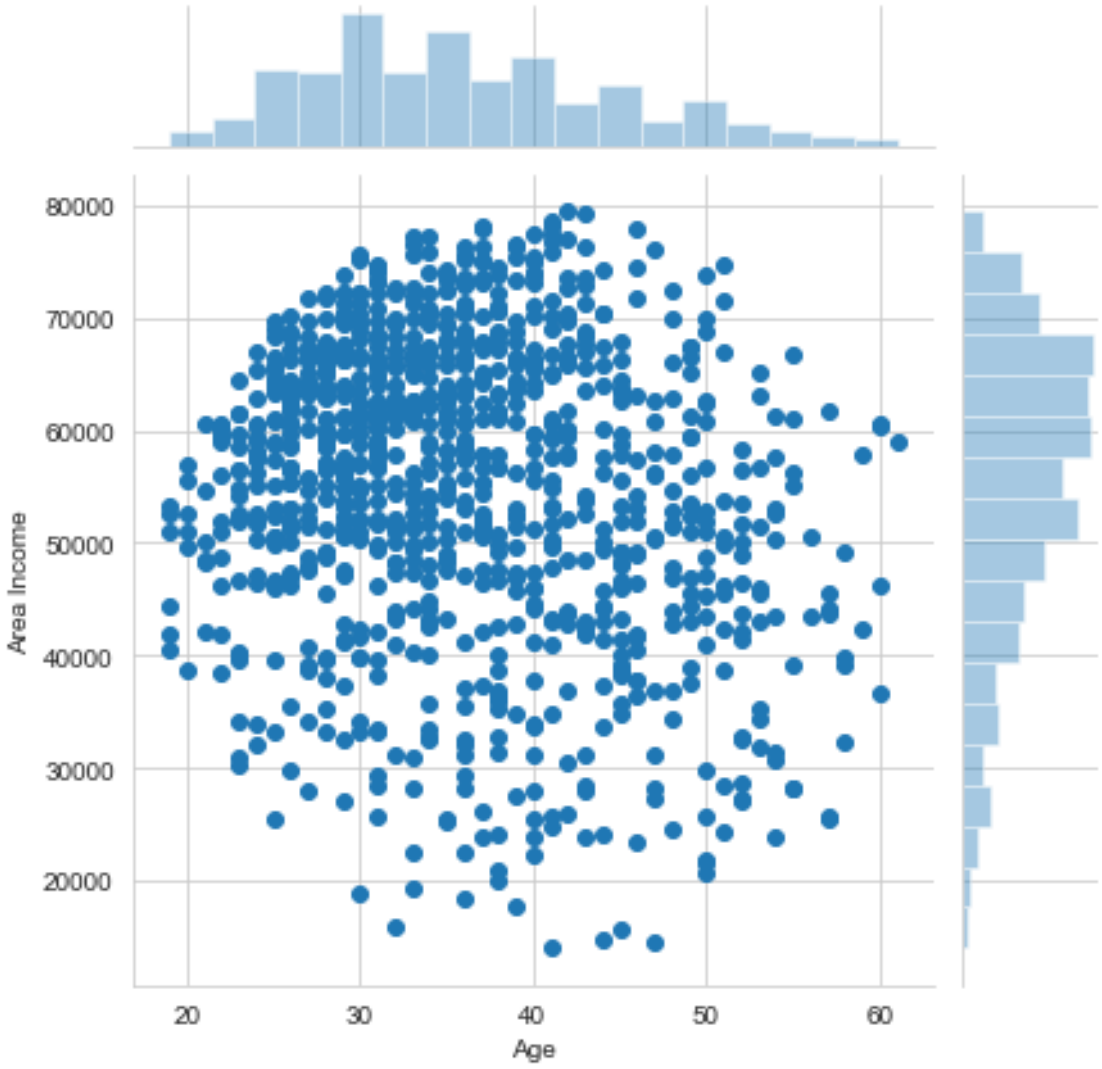

Jointplot showing the kde distributions of Daily Time spent on site vs. Age

[8]:
```
sns.jointplot(x='Age',y='Daily Time Spent on␣
↪Site',data=ad_data,color='red',kind='kde');
```

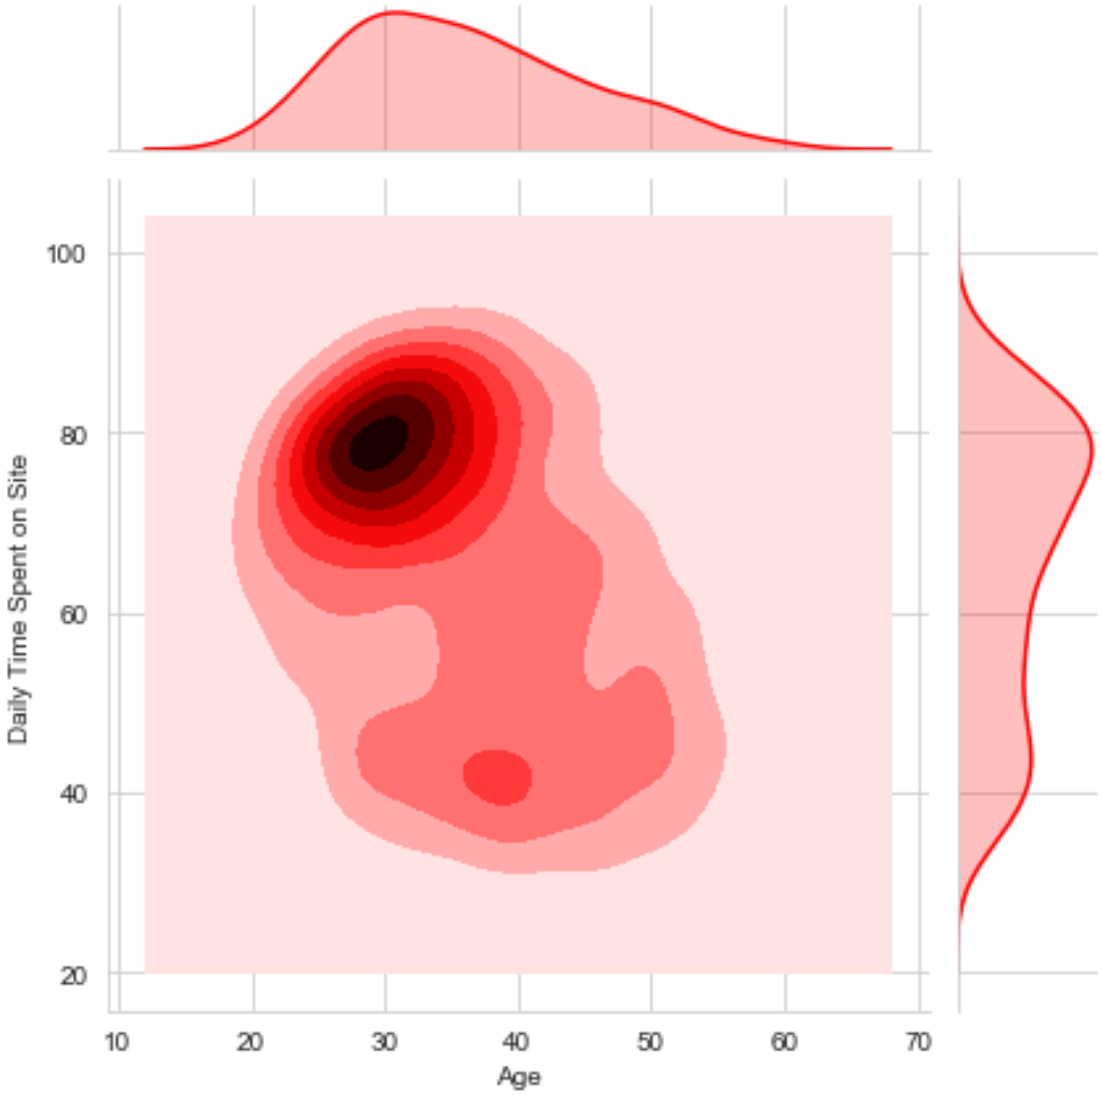

Jointplot of 'Daily Time Spent on Site' vs. 'Daily Internet Usage

[9]:
```
sns.jointplot(x='Daily Time Spent on Site',y='Daily Internet␣
↪Usage',data=ad_data,color='green')
```

[9]:
```
<seaborn.axisgrid.JointGrid at 0x21f5922e388>
```

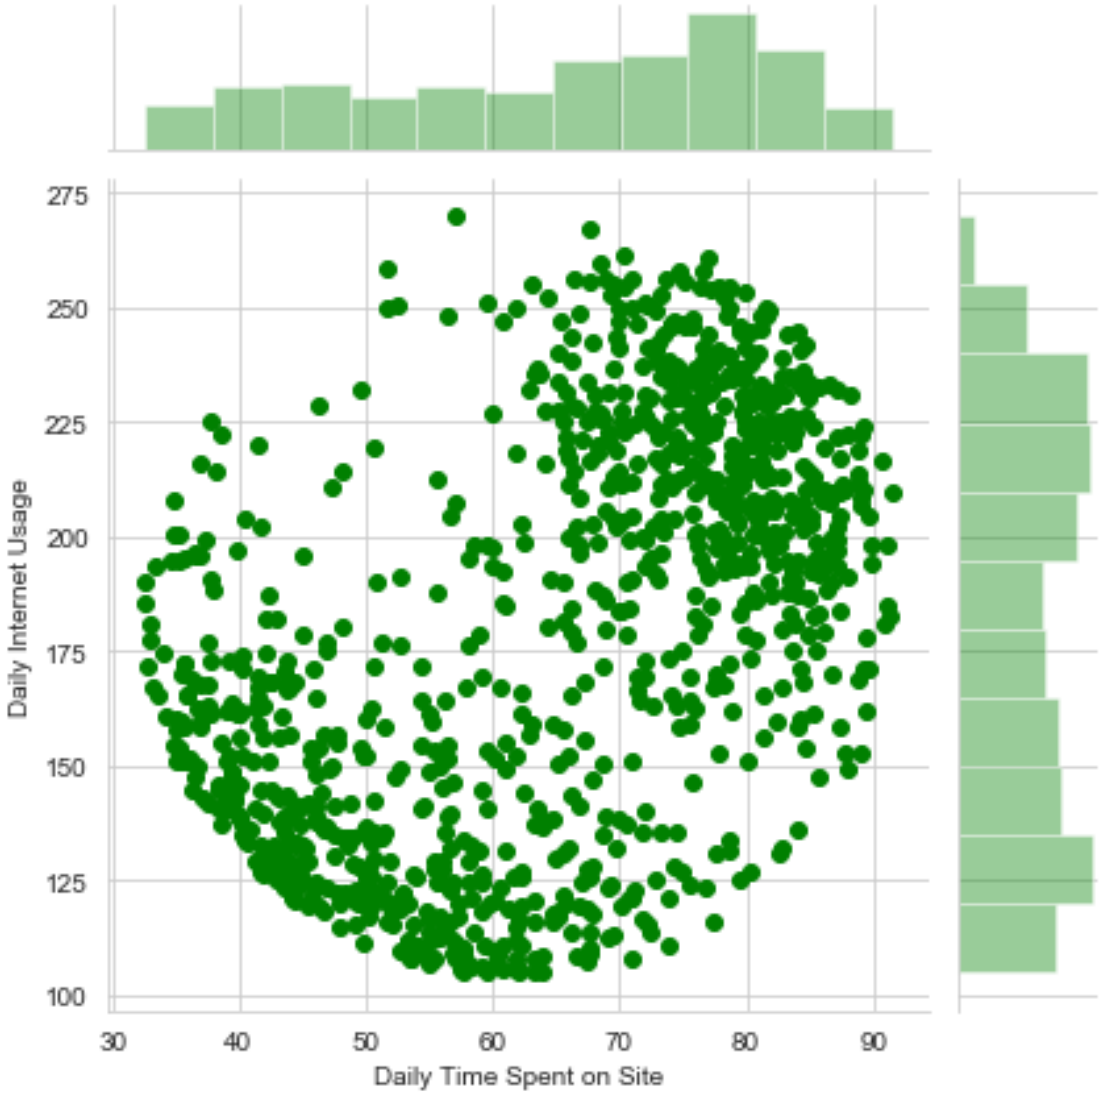

Pairplot with the hue defined by the 'Clicked on Ad'

[10]:
```
sns.pairplot(ad_data,hue='Clicked on Ad',palette='bwr')
```

```
C:\Users\rajar\Anaconda3\lib\site-packages\statsmodels\nonparametric\kde.py:487:
RuntimeWarning: invalid value encountered in true_divide
  binned = fast_linbin(X, a, b, gridsize) / (delta * nobs)
C:\Users\rajar\Anaconda3\lib\site-
packages\statsmodels\nonparametric\kdetools.py:34: RuntimeWarning: invalid value
encountered in double_scalars
  FAC1 = 2*(np.pi*bw/RANGE)**2
```

[10]:
```
<seaborn.axisgrid.PairGrid at 0x21f5921f948>
```

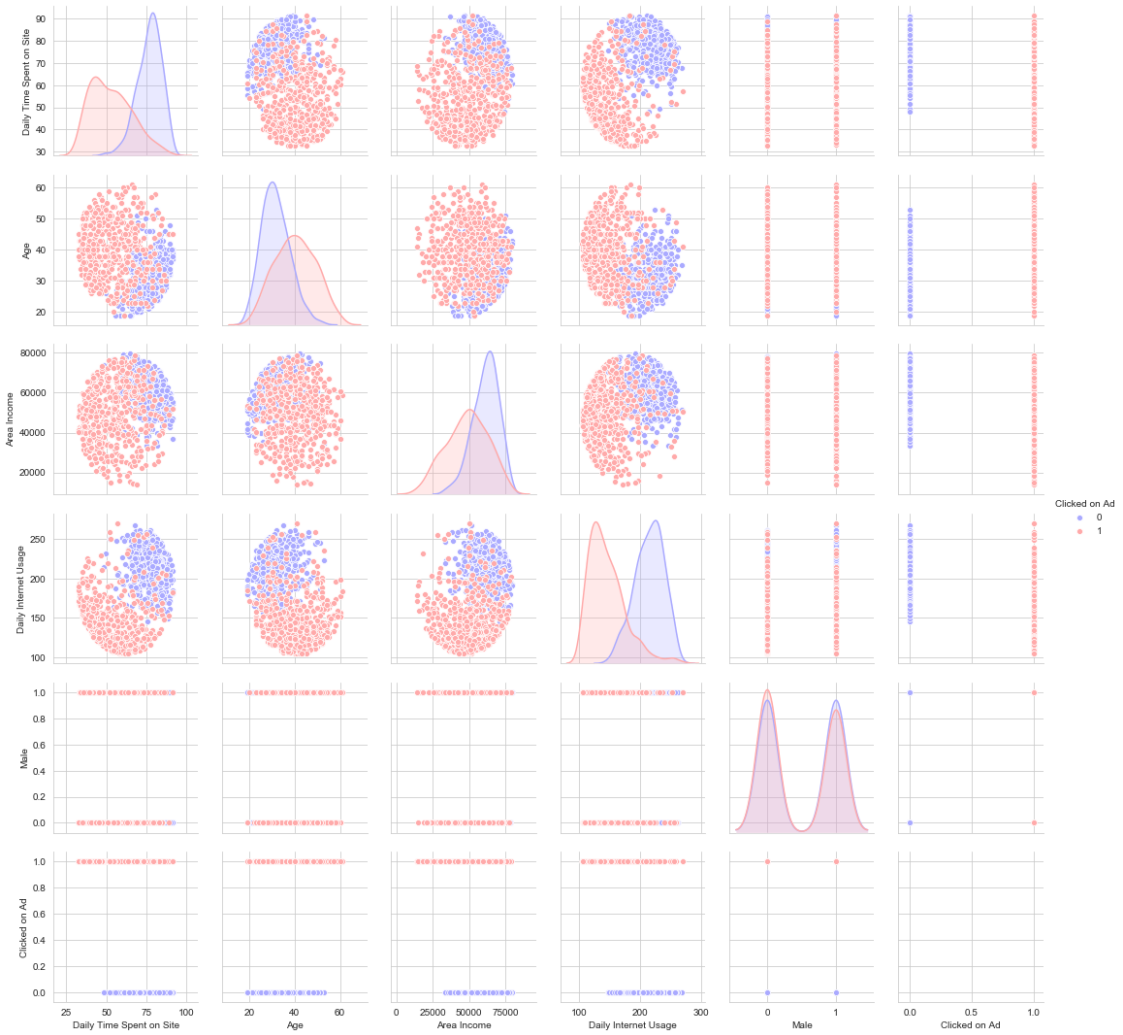

0.0.2 Training and Testing the Data

[11]:
```
from sklearn.model_selection import train_test_split
```

[12]:
```
ad_data.columns
```

[12]:
```
Index(['Daily Time Spent on Site', 'Age', 'Area Income',
       'Daily Internet Usage', 'Ad Topic Line', 'City', 'Male', 'Country',
       'Timestamp', 'Clicked on Ad'],
      dtype='object')
```

[13]:
```
X = ad_data[['Daily Time Spent on Site', 'Age', 'Area Income','Daily Internet␣
 ↪Usage', 'Male']]
y = ad_data['Clicked on Ad']
```

[15]:
```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,␣
↪random_state=42)
```

Training the Model

[16]:
```
from sklearn.linear_model import LogisticRegression
```

[17]:
```
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
```

```
C:\Users\rajar\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432:
FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a
solver to silence this warning.
  FutureWarning)
```

[17]:
```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                   intercept_scaling=1, l1_ratio=None, max_iter=100,
                   multi_class='warn', n_jobs=None, penalty='l2',
                   random_state=None, solver='warn', tol=0.0001, verbose=0,
                   warm_start=False)
```

Predicting the Test Data

[19]:
```
predictions = logmodel.predict(X_test)
```

Evaluating the model

[20]:
```
from sklearn.metrics import classification_report
```

[21]:
```
print(classification_report(y_test,predictions))
```

```
              precision    recall  f1-score   support

           0       0.87      0.96      0.91       162
           1       0.96      0.86      0.91       168

    accuracy                           0.91       330
   macro avg       0.91      0.91      0.91       330
weighted avg       0.91      0.91      0.91       330
```